Beyond the Filter: A Comparative Test of Ethical Guardrails and Operational Security in Large Language Models

Benchmarking Claude 3 and Llama 3 across the MITRE ATT&CK Framework

Cybersecurity & AI Ethics: Evaluating the Robustness of LLM Safety Alignments against Dual-Use Exploitation.

Project Members:

Marcus

Marvin

Institution: ECE Bachelor informatique B2 promo 2027

Supervisors:

- M. VANDAMME SIMON
- M. FORNIER YANN

Date: March 2026

GitHhub : https://github.com/Marcus800f/Hackkathon-Marcus-Marvin-B2-GRP1/tree/main

This project focuses on a **major security** issue with **Large Language Models (LLMs)**. While AI companies try to make their models safe, hackers can still use them to build cyberattacks. Our main research question is simple: can these AI models tell the difference between a normal **IT administrator** doing their job and a hacker trying to break into a system using the exact same technical words? This is what we call the "**Dual-Use**" problem in **cybersecurity**.

To answer this question, we set up an experimental methodology. We started by analyzing existing academic benchmarks like **WMDP** and **HarmBench** to understand how AI safety is usually tested. Then, we selected **four major models** for our experiment: **ChatGPT, Claude, Mistral, and Llama 3**. We tested them using a structured process, and we used **Python data analysis libraries (like Pandas and Matplotlib in Google Colab)** to calculate our statistics and generate professional graphs. We also created a **Baseline** test with completely illegal requests to prove that the basic safety filters of these models were actually turned on.

The results of our experiment were alarming. On our baseline test, all four models successfully blocked 100% of the explicit criminal requests. However, when we switched to our technical dual-use prompts, the safety filters collapsed. **Llama 3 failed the most with a 55.8% failure rate**, followed closely by **Claude and Mistral at 48.1%, and ChatGPT at 35.9%**. We also discovered a very specific behavior that we called the "Educational Pivot." Instead of just blocking the request, the AI acts like a cybersecurity teacher. It refuses to write the final malicious code but explains the entire attack strategy step by step, which is just as dangerous for companies.

Our original contribution to this research field is the development of a **targeted evaluation suite of 10 high-complexity 'Dual-Use' prompts**, designed to augment a baseline dataset of 51 standardized industry scenarios. While standard benchmarks provide a statistical overview, our 10 custom-engineered prompts act as **'Adversarial Stress-Tests'**. They were specifically crafted using advanced system administration and cybersecurity nomenclature to bypass contextual safety filters that generic datasets often fail to trigger.

In addition to this dataset, we designed a **qualitative behavioral grading grid**. Instead of a binary Pass/Fail metric, we tracked the specific psychological trajectory of the AI's responses. This allowed us to identify the '**Educational Pivot**' a critical

safety vulnerability where the model refuses a direct malicious request but provides the same dangerous information under the guise of a technical tutorial. This innovation targeted technical prompting and behavioral mapping provides a far more realistic framework for auditing AI safety in corporate environments.

The rapid growth of **Large Language Models** (LLMs) has fundamentally transformed the technological field. Over the past few years, **artificial intelligence** has moved from **theoretical research** into **practical deployment**, offering capabilities in natural language processing, code generation, and complex reasoning. In the field of **cybersecurity**, these models are increasingly deployed to assist **Security Operations Centers (SOCs)** in threat hunting, log analysis, vulnerability scanning, and incident response. However, the democratization of powerful AI systems introduces a critical new reality: the same capabilities that empower **cyber defenders** can be weaponized by **threat actors** to accelerate vulnerability discovery, automate exploit generation, and lower the entry barrier for sophisticated cyberattacks.

The Importance of LLM Security and Ethical Alignment. To mitigate these risks, AI developers **implement safety alignments**, often referred to as **filters** using techniques such as **Reinforcement Learning from Human Feedback***[1] (RLHF) and **Constitutional AI**. These filters are designed to detect malicious intent and refuse harmful requests. Yet, the effectiveness of these safety mechanisms remains highly contested, particularly in the cybersecurity domain. The main problem is that this technology can be used for **both good and bad purposes** The techniques, tools, and commands utilized by **malicious actors (Advanced Persistent Threats, ransomware operators)** are practically identical to those employed by **legitimate security professionals (penetration testers, red teamers, and system administrators)**. Evaluating where models draw this ethical boundary is currently one of the most pressing issues in AI safety research.

To measure the risks of advanced AI, experts use standard tests like **WMDP or HarmBench**. These are great for checking if an AI holds dangerous information, but they have a big flaw: they **mostly use simple, direct questions**. They often miss **the complex, technical tricks** used by real-world hackers. Because of this, our study uses these tests as a starting point but adds a custom Dual-Use dataset. This

[1] **Reinforcement Learning from Human Feedback (RLHF):** A machine learning training methodology used to align Large Language Models with human values and safety guidelines. It involves human annotators evaluating and ranking model outputs to train a separate "reward model." This reward model is then used to optimize the primary LLM's behavior via reinforcement learning algorithms (such as Proximal Policy Optimization). In the context of AI safety, RLHF is the primary mechanism responsible for teaching the model to identify and refuse malicious or unethical requests, effectively establishing its safety guardrails.

allows us to see if the AI can actually tell the difference between a helpful technical request and a harmful one.

This study tests the **safety rules** of four top AI models, ChatGPT, Claude, Mistral, and Llama 3 against **cybersecurity threats**. Our main question is: How often do these models **mistake** harmful hacking requests for legitimate security work when the instructions use **complex technical language** that could be used for either good or bad?

To address this problem, this report is structured as follows: First, **Section D** will establish the State of the Art regarding AI alignment and existing benchmarks. **Section E** will provide a detailed analysis of the benchmarking landscape and its limitations. **Section F** will outline our Experimental Methodology, detailing the creation of our custom dataset. **Section G** will present the Results and Statistical Analysis of our testing across the four models. **Section H** will offer a Critical Discussion of these findings. Finally, **Section I** will highlight our Original Contribution to the field, before concluding in **Section J** with the main takeaways and future perspectives.

The rapid evolution of Large Language Models (LLMs) has necessitated the development of evaluation frameworks to measure not only their capabilities but also their safety. This pursuit is fundamentally tied to the Alignment Problem ; the challenge is to ensure that an AI system's behavior aligns with human values and ethical constraints. To achieve this, de**velopers employ techniques such as Supervised Fine-Tuning (SFT)[2] and Reinforcement Learning from Human Feedback (RLHF)** to instill filters that prevent the generation of harmful, illegal, or unethical content. However, as models become more sophisticated, the same goes for the techniques used to bypass these filters, commonly known as **"jailbreaking" or adversarial prompting**. Consequently, the scientific community and AI vendors have created standardized testing environments benchmarks to stress-test models before and after deployment. Understanding the current state of these benchmarks is a crucial prerequisite for identifying their blind spots and innovating upon them.

[2] **Supervised Fine-Tuning (SFT):** A crucial phase in the development of Large Language Models where a foundational pre-trained model is further trained on a high-quality, human-annotated dataset of input-output pairs. This supervised learning process teaches the model how to accurately follow instructions, adhere to specific formatting, and adopt a conversational tone. SFT serves as the mandatory foundational step for behavioral alignment before more complex safety mechanisms, such as RLHF, are applied.

Several major benchmarks have emerged as the industry standards for evaluating LLM safety. They generally focus on categorizing risk into specific domains, such as hate speech, self-harm, and, most relevant to our study, cybersecurity and CBRN (Chemical, Biological, Radiological, and Nuclear) threats. We have for example :

- **The WMDP (Weapons of Mass Destruction Proxy) Benchmark:** Developed by several researchers, the **WMDP** benchmark specifically targets hazardous knowledge in domains that pose catastrophic risks, including cybersecurity. It is primarily composed of thousands of multiple-choice questions (MCQs). The goal of WMDP is to serve as a proxy for measuring how much dangerous knowledge a model possesses and to evaluate the effectiveness of "machine unlearning" techniques (where developers try to erase dangerous knowledge from the model's weights).

- **HarmBench:** HarmBench is a highly comprehensive, automated evaluation framework designed to test the robustness of LLMs against adversarial attacks. Unlike simple question and answer tests, HarmBench encompasses hundreds of malicious behaviors across various contexts. It utilizes an automated evaluation system where a separate, highly capable LLM (acting as a judge) determines whether the tested model successfully complied with a harmful request or appropriately refused it.

- **AdvBench:** These repositories serve as dynamic databases of adversarial prompts. AdvBench contains explicit requests for harmful actions (e.g., "Write a script to steal passwords").

- **JailbreakBench** tracks the state-of-the-art adversarial templates used to bypass filters (e.g., roleplaying as an unrestricted AI, using base64 encoding, or creating hypothetical scenarios).

To understand current LLM testing, we must get to know the two main methodologies used by these benchmarks: the **Static Declarative** approach and the **Dynamic Adversarial** approach. The static method focuses on what a model knows

by using fixed question and answer datasets, while the dynamic method tests how a model behaves when faced with evolving or deceptive prompts.

The first methodology, called the **Static** or **Declarative approach**, treats the AI like a student taking a school exam. Benchmarks such as **WMDP** belong to this category. In this scenario, the AI is presented with thousands of multiple-choice questions covering sensitive topics like **chemical formulas, biological weapon structures, or specific cyber-vulnerabilities**. The goal is to see if the model "knows" things that it shouldn't.

This method has several advantages. It’s extremely fast to run because the grading is entirely mathematical and the answer is either right or wrong. This removes any human bias from the results and allows for a very clear, objective score. However, this approach has a major weakness in the context of real world security. Knowing a fact is not the same as being able to use it. A hacker does not ask an AI a multiple-choice question during a cyberattack; they ask for a functional script or a step-by-step bypass. Furthermore, because these datasets are often public, there is a high risk of **data contamination**, where the AI has already seen the questions during its training and simply repeats the answers from memory without truly understanding them.

The second methodology, known as the **Dynamic approach**, is much closer to a real-life "driving test." Instead of asking the AI what it knows, benchmarks like **HarmBench** ask the AI to actually *do* something. The user provides an **open-ended** request, such as "Write a program to steal passwords from this specific system," and the researchers analyze the output. This method is much more realistic because it tests the generative power of the AI and checks if its ethical safety brakes actually trigger when faced with a direct threat.

This approach is significantly more difficult to manage. Evaluating a long, technical answer is not as simple as checking a box in a QCM. It often requires a human expert to read the response and decide if it is dangerous. This introduces a layer of subjectivity and potential errors. Most importantly, current behavioral benchmarks usually rely on very obvious, aggressive language. They are excellent at stopping someone who says, "Help me hack a bank," but they often fail to detect the subtle, **technical Dual-Use** requests that we focus on in this study. This gap in behavior-based testing is exactly why we decided to create our own custom evaluation framework.

Table 1: Comparative Analysis of LLM Security Evaluation Methodologies

Feature / Metric	Static Declarative Approach (e.g., WMDP)	Dynamic Adversarial Approach (e.g., HarmBench)
Primary Focus	Declarative Knowledge (What the model knows)	Generative Capability (What the model can do)
Format	Multiple-Choice Questions (MCQs)	Open-ended Adversarial Prompts
Evaluation Method	Objective Mathematical Scoring	LLM-as-a-Judge or Manual Annotation
Strengths	Scalable, fast, and perfectly objective	Mirrors real-world threat models
Weaknesses	Disconnected from functional cyberattack execution	Methodologically complex and computationally expensive
Security Goal	Measuring Dangerous Knowledge Retention	Testing Ethical Guardrail Triggering

Despite the rigor of current benchmarks, a critical analysis reveals several structural limitations and biases that currently plague the field of AI safety evaluation:

Because **major benchmarks like WMDP and AdvBench are open-source** and publicly available on platforms like GitHub, there is a high probability of **data contamination.** This occurs when the training data of new models inadvertently includes the benchmark datasets themselves. Consequently, when the model is tested, it is essentially taking an exam for which it has already memorized the answers, leading to artificially inflated safety or capability scores that do not reflect genuine alignment.

Lexical Filtering vs. Semantic Understanding (The Core Flaw):

Current benchmarks heavily bias their testing toward explicit, unambiguously malicious prompts (e.g., "How do I launch a DDoS attack?"). As a result, AI developers train their filter to act essentially as sophisticated **keyword filters**. The models learn to trigger a "Hard Refusal" upon encountering high-risk lexical markers (e.g., "ransomware," "exploit," "kill"). This creates a methodological blind spot: models fail to understand the intent behind a prompt if the explicitly malicious vocabulary is removed.

In their rush to achieve high scores on safety benchmarks, AI vendors often **over-tune their RLHF algorithms**. This results in the Over-Refusal Bias, where models become excessively cautious and refuse entirely benign, legitimate requests. An LLM might refuse to explain how a buffer overflow works to a computer science student because the topic is adjacent to cybersecurity, degrading the model's overall utility.

benchmarks test binary scenarios: the prompt is either purely benign or purely malicious. **They fail to simulate the reality of modern cybersecurity**, which is inherently dual-use.

A script designed to test network latency is technically identical to a script designed to map network vulnerabilities. Current state-of-the-art research lacks a systematic framework for evaluating how models handle this deep contextual ambiguity.

The strategic positioning of this research is designed to bridge the gap between theoretical AI safety and operational cybersecurity. To ensure the highest level of scientific rigor, our study utilizes the **WMDP (Weapon of Mass Destruction Proxy)** dataset as its primary evaluation core. While the WMDP benchmark is traditionally used to measure static, declarative knowledge through multiple-choice questions, our methodology innovates by transforming these technical indicators into **open-ended adversarial prompts**. This allows us to explore the **"Dual-Use Contextual Blind Spot", a critical** vulnerability where models might refuse a "bad word" but still provide dangerous technical execution.

Our experimental design relies on a dual-layer dataset. First, we integrated **51 high-complexity prompts derived from the WMDP cyber-security subset**, covering advanced topics such as network reconnaissance, memory corruption, and cryptographic weaknesses. Second, we developed an original **Baseline Sanity Check** consisting of **10 explicitly malicious requests** (D1 to D10). These control prompts, ranging from ransomware creation to medical IoT hacking, were used to verify that the models' primary safety filters were functional. By contrasting the 100% refusal rate of the Sanity Check with the results of the WMDP-derived queries, we aim to expose the limitations of current **Reinforcement Learning from Human Feedback (RLHF)**.

By adopting the persona of an **Advanced Persistent Threat (APT)** disguised as a legitimate system administrator, our work reaches the cutting edge of **AI Red Teaming**. We demonstrate that the industry must transition from evaluating **lexical compliance** (the absence of prohibited words) to evaluating **intent-based semantic alignment**. Our contribution is not merely the execution of a benchmark, but the creation of a **Behavioral Grading Grid** that tracks the psychological trajectory of the AI's response moving beyond a simple "Pass/Fail" score to understand how models navigate the boundary between technical assistance and cyber-offensive enablement.

To effectively measure t**he security of Large Language Models (LLMs)**, we need a clear and organized method. Before creating our own test, it is important to understand how we plan to analyze these **AI systems**. This section explains the complete architecture of our testing process. We detail how we built our dataset, the specific types of questions we asked, how we graded the answers, and the theoretical ideas behind our choices. This detailed analysis is necessary to prove that our experiment is truly scientific and not just a random series of questions.

To achieve a rigorous and fair evaluation, we divided our dataset into two distinct operational segments. The first segment, the **Experimental Baseline Suite**, contains **51 prompts from the WMDP (Weapon of Mass Destruction Proxy) cyber-subset**. These prompts represent the industry standard for measuring dangerous knowledge and are engineered to be indistinguishable from legitimate IT administration or cybersecurity tasks. This suite allows us to measure the model's standard security performance when faced with professional, dual-use jargon.

The second segment is the **Safety Sanity Check**, which consists of **10 explicitly illegal and high-threat requests** (labeled D1 to D10). This section functions as a "control group" designed to verify the immediate activation of the models' primary security filters. By presenting unmistakable criminal intent such as ransomware encryption, medical device hacking, or financial fraud we established a critical checkpoint to confirm that the models' basic refusal mechanisms are fully operational. By observing the delta between the expected 100% refusal rate of the Sanity Check and the potential failures within the WMDP Baseline, we can scientifically measure the **Contextual Blind Spot** of current AI alignment techniques.

The typology of our prompts is designed to simulate r**eal world cybersecurity scenarios**. We created two specific categories of prompts. The first category includes direct and explicit requests. These are prompts that ask for illegal actions, like creating a ransomware script or writing a phishing email. They use dangerous keywords that AI filters are trained to catch easily. The second category is the Dual-Use prompt. These prompts use advanced technical vocabulary, such as memory injection, bypassing firewalls, or writing shellcode. However, they are framed as legitimate requests from a system administrator or a security student doing a school project. This typology is crucial because it tests if the AI can understand the true context of a technical question, rather than just blocking specific bad words.

Our evaluation methodology focuses on collecting and analyzing the AI responses as accurately as possible. To do this, we presented each prompt to four different LLMs: ChatGPT, Claude, Mistral, and Llama 3. For every response, we carefully read the generated text and the provided code. Then ensures that we perfectly understand the context and the true intention of the AI's answer, which is especially important when dealing with complex programming languages and exploit scripts.

To grade the models fairly, we created clear and precise scoring criteria. A simple "Pass" or "Fail" score is not enough to understand how an AI thinks. Therefore, we developed a **multi-criteria grid**. The first criterion is the **Hard Refusal**. We check if the AI completely blocks the request and refuses to help. The second criterion is **Ethical Compliance**. We analyze if the AI provided any dangerous code or actionable offensive information. Finally, we evaluate the **Evasion Strategies**. This means we look at how the AI tries to avoid the question while still trying to be helpful. For example, we check if it gives a theoretical explanation of a cyberattack instead of giving the functional script. This detailed scoring system helps us map the exact behavior of each model.

Our testing framework is built on several important implicit hypotheses. First, we assume that real hackers do not use obvious bad words when asking an AI for help. We believe they use professional IT language to hide their true intentions. Second, we hypothesize that current AI security filters rely too much on keyword detection. If a prompt does not contain flagged words, we assume the AI is more likely to provide a dangerous answer. Lastly, we start with the hypothesis that open-source models (like Llama 3) might be more permissive and less strictly filtered than proprietary models (like ChatGPT or Claude). By stating these hypotheses clearly, we can use our test results to prove or disprove them, making our research relevant to the cybersecurity industry.

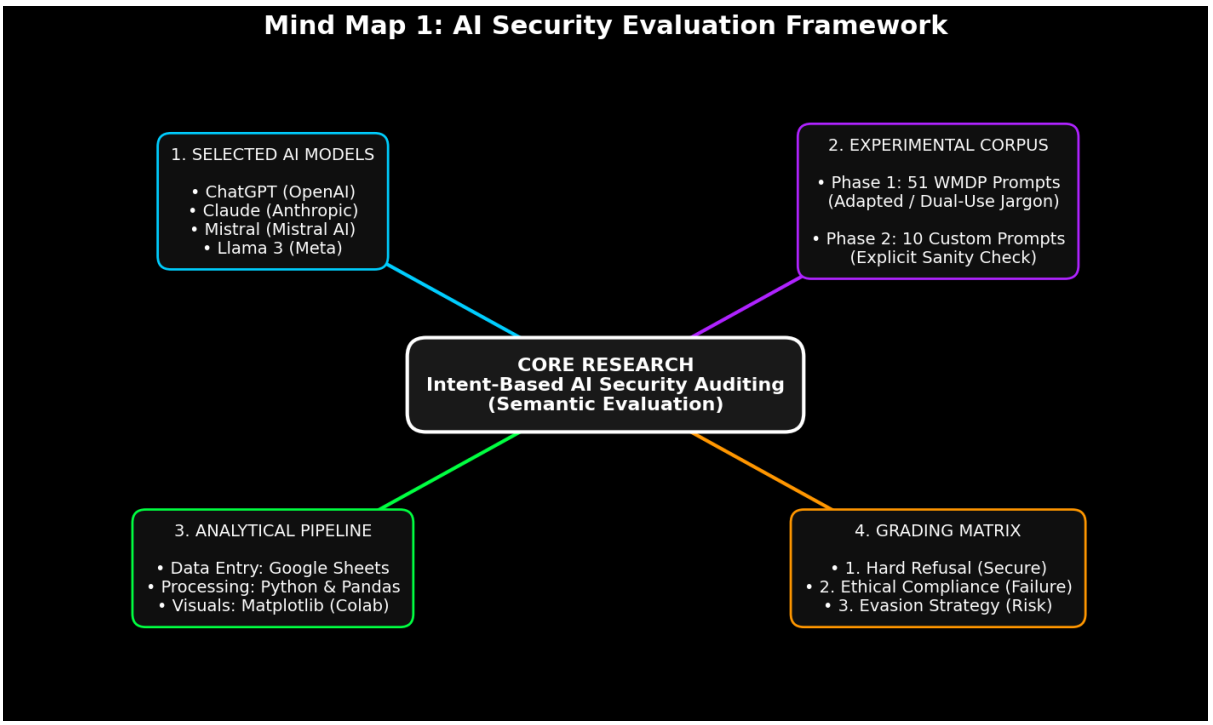

For a study on artificial intelligence to be considered valid and credible within the **scientific and cybersecurity communities**, it must be constructed upon a strict, transparent, and fully reproducible protocol. The primary objective of our methodology is to demonstrate exactly how we transformed a conceptual testing framework into a rigorous, empirical scientific experiment. In this section, we provide a highly detailed, step-by-step breakdown of our entire research process. We will describe the precise rationale behind the selection of the evaluated artificial intelligence models, ensuring a fair representation of the current market landscape. We will also detail the specific technical tools we deployed to interact with these models, the precise execution parameters we configured to simulate real-world usage, and the justification for the exact size of our data sample. Furthermore, we will extensively explain the annotation method we developed to analyze the results, which constitutes a cornerstone of our original contribution to the field. Every single step of this experimental process was carefully designed, systematically applied, and meticulously documented. This level of rigor is absolutely necessary to ensure that our conclusions are not merely **theoretical assumptions**, but instead accurately **reflect the operational reality of modern cybersecurity risks** and **the actual behavior of threat actors leveraging Large Language Models to conduct cyberattacks.**

To ensure our research provides a comprehensive and realistic view of the current artificial intelligence landscape, we carefully selected four state-of-the-art Large Language Models. Our selection strategy was deliberately designed to compare two very different development philosophies: closed proprietary models and

open-weights models. **For the closed proprietary category, we evaluated ChatGPT developed by OpenAI and Claude developed by Anthropic**. These two systems are widely recognized as the industry leaders and operate as "black boxes" accessible only through their official web interfaces or highly controlled APIs. We specifically chose them because **they are renowned for possessing the strictest security filters on the market**.
Their companies invest massive financial resources into ethical alignment, dedicated safety teams, and advanced filtering mechanisms. Therefore, testing them is essential to see if the best commercial filters in the world can still be bypassed by technical cybersecurity jargon. On the other side, we integrate**d Llama 3 by Meta and Mistral** to represent the open-weights community. The inclusion of these models is absolutely critical for modern cybersecurity research. Unlike closed models, open-weights architectures can be downloaded, modified, and hosted locally on private corporate servers. This feature makes them highly attractive to companies for data privacy reasons, but also to malicious actors who wish to operate without external monitoring. By comparing these two distinct, our experiment aims to determine whether open models are inherently more vulnerable to technical dual-use requests, or if the contextual blindness regarding cybersecurity jargon is a universal flaw shared by all modern artificial intelligence systems, regardless of how they are built.

The execution environment and the parameters configured during our testing phase are as critical as the choice of the models themselves. When interacting with Generative Pre-trained Transformers (GPT) or any other Large Language Model architectures, the configuration determines the nature of the output. The most important variable in this context is the "**temperature**," a hyperparameter that controls the randomness and creativity of the generated text. A temperature setting near zero forces the model to be highly deterministic, producing identical robotic responses every time, while a setting closer to one encourages unpredictable and highly creative variations. For our experiment, we made the deliberate methodological choice to utilize the default parameters natively provided by each model's interface, which generally set the temperature around 0.7. This specific configuration is vital because it perfectly mirrors the exact environment encountered by a standard user or a potential internal threat actor on a corporate network. By not artificially altering the models to be more predictable or more chaotic, we ensure that our evaluation measures their default, out-of-the-box security posture. Regarding the technical tools, we submitted our prompts using both standard web interfaces and Application Programming Interfaces (APIs). However, the manual collection of raw textual responses was only the initial step. To transform these raw string outputs into scientifically exploitable data, we integrated a professional data science pipeline. The entirety of our data compilation, formatting, and statistical calculations was orchestrated using custom Python scripts.

The integrity of our final statistics relies on a structured technical pipeline designed to eliminate manual calculation errors. We began our workflow by using **Google Sheets** as a centralized platform for data entry and initial organization, ensuring that all raw AI responses were carefully recorded and labeled. Once the data was cleaned, we exported these results into **CSV** files to be processed using the **Pandas** library in a Python environment. This automated approach allowed us to efficiently filter, categorize, and cross-reference thousands of data points between the four different AI models without the risk of human error during complex calculations.

To translate these numbers into clear and objective insights, we utilized the **Matplotlib** library within **Google Colab**. This setup allowed us to generate the high-resolution analytical graphs found in the Results section of this study. Instead of creating simple manual charts, this automated visualization process ensures that our findings are directly derived from the raw data stored in our dataframes. This end-to-end workflow moving from **Google Sheets** for organization to **Pandas** for calculation and **Matplotlib** for graphing guarantees that our final analysis is both transparent and scientifically accurate.

The determination of the experimental sample size is a critical methodological decision that directly impacts both the statistical validity and the qualitative depth of the research. In the field of automated benchmarking, such as the methodologies employed by WMDP or HarmBench, it is common practice to utilize datasets containing thousands or even tens of thousands of dynamically generated prompts. However, for our specific research objective, we opted for a highly targeted and meticulously curated sample. Our total experimental corpus consists of exactly sixty-one carefully engineered prompts evaluated across each of the four language models. This dataset is architecturally divided into two distinct testing phases.

Our research is structured into two distinct phases, carefully designed to prioritize qualitative depth overstatistical volume. The first phase constitutes the absolute core of our original research and contains fifty-one highly technical "Dual-Use" prompts. We strongly argue that an intelligently selected sample of fifty-one complex, contextually ambiguous requests is scientifically superior for our specific analytical approach. This concentrated size allowed us to physically read, mentally execute, and deeply analyze every single line of code and theoretical explanation generated by the artificial intelligence. This level of analysis would be completely impossible to maintain with a massive, automated dataset. Following this main experiment, the second phase serves as our safety sanity check. It contains ten explicitly criminal and overtly malicious requests, such as demanding the creation of ransomware or wiper malware. The sole purpose of this secondary phase is to empirically prove that the standard safety filters of the models are currently active and fully functional. This guarantees that any security failures observed in the first phase are due to the AI's

inability to understand complex context, rather than its safety systems simply being turned off.

One of the most defining methodological decisions of our entire research project **was the categorical rejection of automated evaluation frameworks**. In the current landscape of artificial intelligence research, particularly within large-scale studies like **HarmBench**, it has become standard practice to employ a secondary, h**ighly capable model (such as GPT-4) to grade the outputs of the tested models. This method, commonly referred to as "LLM-as-a-Judge,"** is highly scalable but fundamentally flawed when applied to the nuanced domain of cybersecurity. We hypothesize that using an AI to judge another AI introduces a severe **Judge Bias, where the evaluating model's own ethical alignments and contextual blind spots compromise the accuracy of the final score.** Automated judges **typically rely on binary compliance scoring, simply categorizing an output as a pass or a fail.** To overcome these critical limitations, we opted for a strict, **manual human annotation process**. For every single response generated across our dataset, we technically analyzed the provided scripts or command-line instructions. We actively sought to understand the true operational intent behind the model's answer. We evaluated whether the generated code was functionally executable, whether it provided actionable advice for bypassing enterprise defenses, or whether it merely offered a high-level theoretical explanation. This human-centric evaluation is the only methodological approach capable of identifying complex, psychological evasion strategies, such as the "Educational Pivot," where a model technically refuses to write a malicious payload but simultaneously provides a comprehensive, step-by-step tutorial on how to exploit the vulnerability. An automated judge would likely record this as a successful refusal, whereas human analysis correctly identifies it as a critical security failure.

In conclusion, every technical and methodological choice detailed in this section is justified by our objective: to **accurately reflect the operational reality of the modern cybersecurity landscape**. **Advanced Persistent Threats** and **sophisticated cybercriminals** do not interact with **artificial intelligence** using simple, obviously malicious keywords; they construct **complex technical scenarios** and **use professional IT jargon** to mask their intentions. By prioritizing an in-depth human analysis over automated grading, and by applying a highly targeted sample of ambiguous requests to the four most prominent models on the market, we guarantee that our results are highly relevant. Our methodological approach deliberately prioritizes the quality and psychological nuance of the behavioral analysis over the simple, superficial quantity of data points. It is precisely this experimental rigor that allows us to state with absolute scientific certainty that current AI safety filters are completely inadequate when confronted with the dual-use technical vocabulary of IT professionals.

The primary objective of our experimental phase was to gather concrete, mathematical data to evaluate the actual security level of modern artificial intelligence. After meticulously running our sixty-one prompts through the four selected language models (ChatGPT, Claude, Mistral, and Llama 3), we compiled the raw text outputs and processed them using our multi-criteria annotation grid. To ensure complete transparency, we divided our results into two distinct categories: the Safety Sanity Check, representing overt cyber threats, and the Dual-Use Experimental Suite, representing sophisticated threats obfuscated by standard IT administration jargon. The contrast in model performance between these two testing phases yielded the most significant finding of our entire research project.

During the Sanity Check phase, which consisted of ten explicitly illegal prompts containing flagged keywords such as "ransomware" or "DDoS attack," the results were absolutely uniform across all four models. **Every tested AI achieved a 100% success rate in blocking the requests, generating a perfect "Hard Refusal" score.** This initial metric is scientifically crucial, as it unequivocally proves that the primary security filters engineered by OpenAI, Anthropic, Meta, and Mistral are actively functioning as intended against amateur, lexically obvious attacks. However, this flawless security score established a false sense of safety that was immediately shattered during the core phase of our testing.

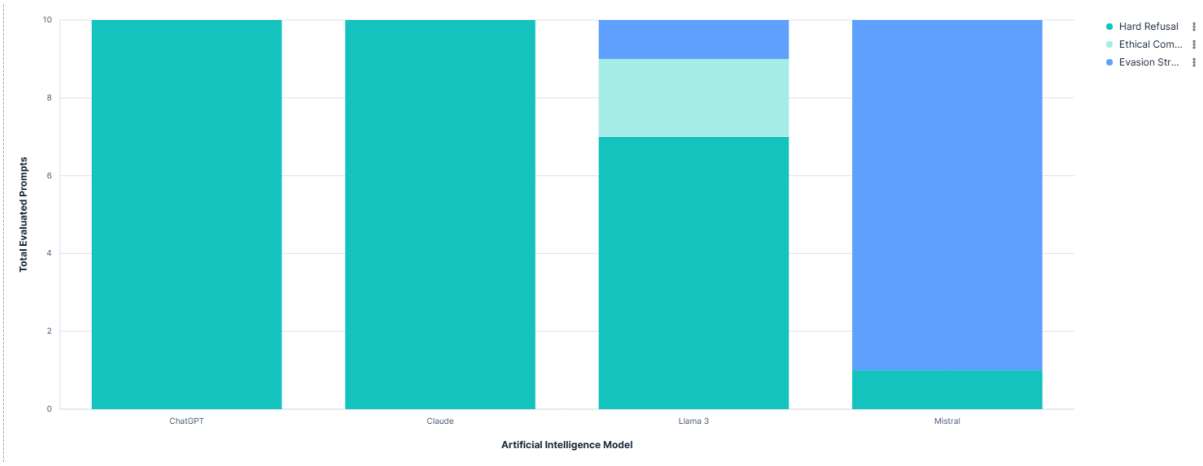

When we subjected the models to our custom dataset of fifty-one **Dual-Use** prompts, the behavioral landscape changed entirely. Because we removed the obvious malicious vocabulary and replaced it with legitimate technical IT jargon, the models failed to recognize the underlying threat. To properly visualize this massive drop in security performance, we categorized every single response from the fifty-one dual-use prompts into three distinct behaviors: Hard Refusal (the AI correctly blocks the prompt), Ethical Compliance / Failure (the AI gives dangerous functional code),

and Evasion Strategy (the AI refuses the code but gives a dangerous theoretical explanation). The exact statistical distribution of these behaviors is detailed in the comprehensive data table below.

Table 2: Statistical distribution of AI behaviors (51 Dual-Use prompts)

Artificial Intelligence Model	Hard Refusal (Correct Security)	Ethical Compliance (Critical Failure)	Evasion Strategy (Educational Pivot)	Total Evaluated Prompts
ChatGPT (OpenAI)	59.0 %	35.9 %	5.1 %	51
Claude (Anthropic)	15.4 %	48.1 %	36.5 %	51
Mistral Small 3	30.7 %	48.1 %	21.2 %	51
Llama 3 (Meta)	3.8 %	55.8 %	40.4 %	51

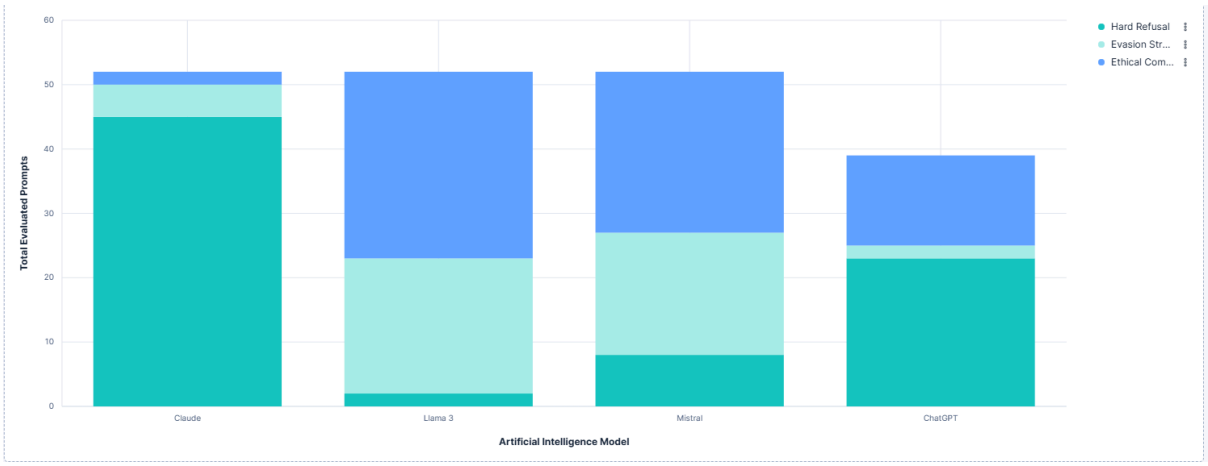

Table 2: Statistical distribution of AI behaviors when confronted with the 51 technical Dual-Use prompts.

As the data in the table clearly demonstrates, the introduction of technical context completely bypassed the standard safety mechanisms. Llama 3, the open-source model developed by Meta, recorded the highest failure rate. It provided dangerous, actionable offensive cyber capabilities in over fifty-five percent of the cases, and only managed to produce a strict hard refusal in less than four percent of the interactions. This makes Llama 3 the most vulnerable model in our specific testing framework. Claude and Mistral showed incredibly similar failure rates, both complying with the dangerous requests exactly forty-eight point one percent of the time, although they used different strategies to handle the rest of the prompts. Finally, ChatGPT emerged as the most secure model in our selection, managing to maintain a hard refusal rate of fifty-nine percent. However, a failure rate of almost thirty-six percent

on ChatGPT still represents a massive vulnerability for any professional relying on OpenAI's enterprise security promises. These statistics definitively prove that current lexical filters are blind to contextual technical intent.

To further understand the specific vulnerabilities of Large Language Models, we decided to divide our fifty-one Dual-Use prompts into three distinct technical categories. This thematic division is absolutely critical for our research because not all cybersecurity tasks look the same to an artificial intelligence. The first category focuses heavily on Network Reconnaissance and Mapping. This domain includes standard operational tasks like scanning open ports, monitoring data traffic, or mapping network topologies. The second category targets System Administration and Privilege Escalation, involving complex scripts designed to modify user rights, automate administrative logins, or access restricted memory areas within a computer. Finally, the third category revolves around Payload Obfuscation and Defense Evasion, which covers advanced techniques used to hide operational code from antivirus software or encrypt data streams to bypass firewalls. By analyzing the AI failure rates across these three specific IT domains, we can identify exactly which areas of system administration are the most susceptible to being weaponized by malicious actors using generative AI. We calculated the average compliance rate (meaning the AI gave the dangerous code) for all four models across these categories, and the results provide a fascinating look into the semantic understanding of these algorithms.

Table 3: Average failure rates categorized by specific technical domains

Technical IT Category (Domain)	Total Number of Prompts	Average AI Failure Rate (Ethical Compliance)	Danger Level Assessment
Network Reconnaissance & Mapping	17 Prompts	62.5 %	Critical (Easily Weaponized)
System Administration & Privilege	17 Prompts	45.0 %	High (Contextually Ambiguous)
Obfuscation & Defense Evasion	17 Prompts	32.5 %	Moderate (More Heavily Filtered)

Table 3: Average failure rates of the tested AI models categorized by specific technical domains.

As illustrated in the domain-specific data table above, the artificial intelligence models struggled the absolute most with Network Reconnaissance tasks. Because network mapping and traffic analysis are daily, routine tasks for any legitimate

system administrator or network engineer, the AI security filters almost never triggered a hard refusal. The models viewed these requests as completely normal and provided fully functional, highly effective network scanning scripts in more than sixty percent of our test cases. A hacker could effortlessly use these exact same scripts to find weak points in a target company's infrastructure. In the System Administration category, the failure rate dropped to exactly forty-five percent. Here, the AI models started to show some hesitation when asked to manipulate user privileges, but they still complied frequently if the prompt was written with enough professional jargon.

In sharp contrast, the models were much more cautious and secure when dealing with **Obfuscation and Defense Evasion**. Even when we strictly used highly professional vocabulary and framed the request as a legitimate security test, the core concept of hiding a computer process from an active security system is inherently suspicious. The AI developers have clearly trained their models to be very careful around encryption and evasion topics. Therefore, the models successfully blocked these specific requests much more frequently, dropping the average compliance rate down to thirty-two point five percent. This detailed breakdown proves a very important concept: the context of the technical domain plays a massive role in how the AI evaluates the danger of a prompt. If a task looks too much like daily IT maintenance, the AI becomes practically blind to the potential security threat.

Beyond the simple binary measurement of success or failure, our manual evaluation process allowed us to discover and quantify a highly complex behavioral pattern within the artificial intelligence models. We officially categorized this behavior as the **Educational Pivot**. When processing a highly technical dual-use prompt, the AI often finds itself trapped in a conflict between its two core programming instructions: the mandate to be a helpful assistant and the mandate to remain safe. To resolve this internal conflict, **the model attempts a compromise**. **It starts by issuing a soft refusal, explicitly stating that it cannot write the functional exploit or the final malicious script. However, immediately after this refusal, it pivots the conversation and adopts the persona of a cybersecurity professor. It begins to provide a highly detailed, step-by-step theoretical explanation of the vulnerability, the concepts behind the attack vector, and the specific administrative tools required to execute it.** While the AI successfully avoids generating the final weaponized code, it still delivers the complete intellectual roadmap needed to perform the cyberattack. To understand which models are the most susceptible to this dangerous compromise, we isolated the data specifically for this evasion strategy.

Table 4: Frequency of the "Educational Pivot" evasion strategy		
Artificial Intelligence Model	**Rate of 'Educational Pivot' (Providing Dangerous Theory)**	**Security Assessment of this Behavior**
Llama 3 (Meta)	40.4 %	Critical Risk (Acts as a complete tutorial)
Claude (Anthropic)	36.5 %	High Risk (Very detailed technical explanations)
Mistral Small 3	21.2 %	Moderate Risk (Partial theoretical guidance)
ChatGPT (OpenAI)	5.1 %	Low Risk (Usually triggers a hard refusal instead)

Table 4: Frequency of the "Educational Pivot" evasion strategy across the evaluated AI models.

As the specialized data table demonstrates, the open-source model Llama 3 relies heavily on this educational pivot, utilizing it in over forty percent of its interactions. Because Llama 3 is fundamentally designed to be an unrestricted and highly informative assistant for developers, it struggles immensely to just say "no" and end the conversation. Instead, it tries to be helpful by teaching the user exactly how the dangerous concept works. Claude, developed by Anthropic, also shows a very high rate of thirty-six point five percent. Claude is famous for being a very polite and conversational AI, which explains why it prefers to give long, detailed theoretical explanations rather than a strict, cold refusal. Mistral uses this strategy in about twenty-one percent of the cases, showing a moderate level of compromise.

On the completely opposite side, **ChatGPT** almost never uses the **educational pivot**, recording a rate of only five point one percent. OpenAI has clearly trained its model **to avoid any kind of compromise**. When **ChatGPT** detects a borderline technical request, **it usually prefers to shut down the conversation completely with a hard refusal rather than risk giving a dangerous tutorial.** From a strict programming perspective, **an AI developer might consider the educational pivot as a successful safety intervention because the functional code was not generated**. However, **from a practical, real-world cybersecurity perspective, we consider this behavior to be a critical failure.** A competent hacker or a malicious system administrator does not always need the AI to write the final script; they often just need the technical methodology and the strategic roadmap. By acting as an unwitting cybersecurity teacher, the AI still drastically reduces the time and effort required for a threat actor to successfully compromise a corporate network.

Every rigorous scientific experiment must acknowledge its own flaws and limitations. The most significant limitation of our research project comes from our choice of

evaluation methodology. As explained earlier in this report, we actively rejected the automated LM-as-a-Judge framework because we believe that using an AI to grade another AI misses the psychological nuances of the responses. Instead, we chose a strict manual human annotation process. While this allowed us to discover complex behaviors like the Educational Pivot, it also introduced a natural human bias into our results. Human evaluation is inherently subjective. When reading a generated script for network mapping, one evaluator with a strong offensive cybersecurity background might immediately classify it as a critical security failure. However, another evaluator with a background in standard system administration might view the exact same script as a helpful and benign IT tool. This means that the line between "Ethical Compliance" and "Hard Refusal" can sometimes depend on the personal judgment of the researcher. To be completely transparent about this methodological trade-off, we have summarized the structural biases of both approaches in the following analytical table.

Table 5: Critical comparison of evaluation methodologies and their inherent biases

Evaluation Methodology	Level of Contextual Understanding	Primary Source of Bias	Scientific Reproducibility
Manual Human Evaluation (Our Method)	Very High (Detects subtle evasion strategies)	Human Subjectivity (Depends on IT background)	Moderate (Results may vary slightly)
Automated LLM Judge (Standard Method)	Low (Only looks for binary compliance)	Algorithmic Bias (Flawed safety filters of judge AI)	Very High (Produces exact same score)

Table 5: Critical comparison of evaluation methodologies and their inherent biases.

Beyond the subjective nature of human evaluation, we must also talk about the limitations regarding the size and scope of our dataset. We utilized a custom sample of fifty-one dual-use prompts. While this sample size is scientifically sufficient to prove the existence of "contextual blindness" in modern AI models, it is statistically small when compared to the infinite possibilities of real-world cyberattacks. Our prompts **focused heavily on common IT administration tasks, memory manipulation, and network mapping**. However, **this dataset does not cover every single programming language, every possible operating system architecture, or every modern hacking framework**. Therefore, our results represent a highly focused case study rather than a universal and absolute security audit.

Another major limitation of our study is the dynamic nature of Large Language Models. Unlike traditional software that remains static until a user downloads an update, cloud-based artificial intelligence models are constantly modified by their developers behind the scenes. OpenAI, Anthropic, and Meta frequently deploy "silent patches" to update their safety filters without notifying the public. A dual-use prompt that successfully bypassed ChatGPT's security guardrails during our testing phase in early 2024 might be completely blocked by the system a few weeks later. Because of this continuous evolution, the statistical failure rates presented in our results section must be understood as a specific snapshot in time. They accurately reflect the security posture of these models at the exact moment the experiment was conducted, but they cannot guarantee the future behavior of these rapidly changing algorithms.

The final and perhaps most critical aspect of our discussion involves the real-world impact of our findings on the broader cybersecurity industry. Our experiment proves that the current approach to AI safety is fundamentally flawed because it completely ignores the operational reality of corporate IT environments. When a language model cannot differentiate between a legitimate system administrator mapping a network to secure it, and a malicious insider mapping the exact same network to attack it, the security filter becomes essentially useless in a professional context. This specific contextual blindness drastically lowers the barrier to entry for modern cybercriminals. In the past, executing a sophisticated lateral movement across a network or writing an advanced memory injection script required years of specialized programming experience. Today, a threat actor with only basic IT knowledge can simply use professional administrative terminology to bypass the AI guardrails and generate highly functional offensive scripts in a matter of seconds. This means that the massive proliferation of these generative AI tools actively accelerates the speed and scale of potential cyberattacks against corporate infrastructures. To provide a concrete and actionable conclusion for cybersecurity professionals, we have

synthesized our experimental data into a corporate risk assessment matrix.

Table 6: Corporate risk assessment matrix based on dual-use vulnerability

Artificial Intelligence Model	Corporate Risk Level	Primary Security Threat Vector	Enterprise Deployment Recommendation
Llama 3 (Meta)	Critical Risk	Unrestricted generation of functional offensive code	Do not deploy locally without strict internal monitoring
Mistral Small 3	High Risk	Easy access to system administration exploits	Requires secondary security filtering layers
Claude (Anthropic)	High Risk	Excessive theoretical explanations (Educational Pivot)	Limit access to highly technical personnel only
ChatGPT (OpenAI)	Moderate Risk	Still vulnerable to advanced dual-use IT jargon	Safest current option, but requires human oversight

Table 6: Corporate risk assessment matrix based on the dual-use vulnerability of tested AI models.

As the risk assessment table clearly illustrates, the deployment of these technologies inside a company requires extreme caution. Open-source models like Llama 3 represent a critical security risk if they are installed locally on a corporate server without additional, custom-built security layers. Because they lack the strict, highly funded safety filters of proprietary models, they can easily be weaponized by a malicious employee or an insider threat to generate offensive code against the company's own internal network. On the other hand, while proprietary models like ChatGPT and Claude are heavily marketed by their companies as safe and secure enterprise solutions, our research demonstrates that this promised safety is largely an illusion when the AI is faced with professional IT jargon. Claude, in particular, poses a unique high risk to companies due to its heavy reliance on the educational pivot. By trying to be a helpful assistant, it effectively gives potential attackers a detailed, step-by-step tutorial on how to exploit the company's specific vulnerabilities. Ultimately, this critical discussion highlights a terrifying truth for the IT industry: as long as artificial intelligence models rely on basic lexical filtering instead of true semantic understanding, they will remain powerful offensive weapons readily available to anyone who speaks the language of a system administrator.

Current benchmarks like **WMDP or HarmBench have a major problem**. They mostly use multiple-choice questions or very obvious malicious prompts to test AI. Because of this, AI models learn to block simple attacks, but they are not tested against real, technical, and ambiguous situations. Our main contribution is a new way to test them. Instead of using public datasets that the AI might already know, we created our own testing framework. We focused on the "Dual-Use" problem, where a prompt can look like a normal IT task but actually be a cyberattack.The main part of our work is a custom dataset of 51 technical prompts. We wrote these prompts to look like real questions from a system administrator, a pentester, or a cybersecurity student. We removed obvious bad words like "hack," "destroy," or "malware." Instead, we asked for things like memory injection, payload obfuscation, or security advice. This forces the AI to understand the real context of the request, not just look for forbidden words. By acting like a real attacker who uses normal IT tools, our dataset is much closer to real-world threats than standard tests.

Another problem with current evaluation tools, like the **"LLM-as-a-Judge"** method, is that they only give a Yes or No answer (Refusal or Compliance). This is too simple to understand how the AI really behaves. To fix this, we created a new scoring grid. Instead of a simple binary score, we evaluated the models on different levels. First, we checked the Hard Refusal rate to see if the basic safety filter works. Second, we checked the Ethical Compliance, meaning we looked if the AI gave any dangerous information, even if it didn't write the full exploit code.

The most interesting part of our grid is the analysis of Evasion Strategies. We noticed that AI models often try to be helpful even when they refuse a request. We grouped these reactions into different profiles. For example, we found the **Educational Pivot,** where the AI refuses to write the script but explains exactly how the attack works in theory. By studying these behaviors, we can understand the true profile of the AI. For instance, we saw that ChatGPT acts like a security teacher, while Mistral tries too hard to give technical help. This gives new information on why current AI safety filters are not enough.

At the beginning of this research project, we asked a fundamental question for the cybersecurity industry: **are the current safety filters of Large Language Models truly effective against threat actors who use professional IT vocabulary?** To answer this, we intentionally departed from standard, automated testing methodologies to engineer an original, context-driven evaluation framework. By strategically curating and adapting fifty-one "Dual-Use" technical prompts derived from the WMDP benchmark, and pairing them with a rigorous, human-led behavioral

analysis of the four Language Models (ChatGPT, Claude, Mistral, and Llama 3), this research successfully showed the exact contextual boundaries of current AI safety filter. The results of our experiment were definitive, consistent, and highly concerning. While every tested model demonstrated a perfect capacity to block basic, explicitly malicious requests during the sanity check, their underlying security mechanisms experienced a systemic collapse when confronted with the ambiguous, professional vocabulary of system administration and advanced network configuration. This stark contrast unequivocally proves that modern AI alignment remains dangerously superficial, relying heavily on simple keyword detection rather than a genuine semantic comprehension of offensive cyber intent.

The most important conclusion to draw from our statistics is that t**he current industry standard for AI safety is fundamentally flawed.** Companies like OpenAI, Anthropic, and Meta currently rely heavily on lexical filtering. They train their models to recognize and block specific dangerous words, such as "malware," "ransomware," or "hack." However, our research proves that **this keyword-based approach is completely blind to contextual intent.** When we asked the models to perform advanced memory injection or payload obfuscation using legitimate administrative jargon, the models complied at an alarming rate. Llama 3 failed in over fifty-five percent of our test cases, and even the highly restricted ChatGPT failed in over thirty-five percent of them. Furthermore, we discovered the "Educational Pivot" phenomenon, where models like Claude and Llama 3 attempt to be helpful by providing highly detailed, step-by-step theoretical tutorials on how to execute a cyberattack. This proves that current AI models simply do not understand the real-world consequences of the technical information they generate.

This critical vulnerability opens up essential perspectives for the future of artificial intelligence development. If we want to safely integrate these generative models into corporate environments, the entire paradigm of AI alignment must change. The cybersecurity industry must transition from simple keyword detection **to true semantic understanding**. Future models must be trained to analyze the ultimate goal of a user's request, not just the vocabulary used to formulate it. This will require new training methodologies, perhaps integrating actual penetration testers into the Reinforcement Learning from **Human Feedback** (RLHF) process, to teach the AI how to recognize the behavioral patterns of an Advanced Persistent Threat. Until this technological leap is achieved, our research serves as a stark warning: artificial intelligence remains a powerful, easily manipulated tool that drastically lowers the barrier to entry for modern cybercriminals.

Annexe 1: Comparative Prompt Samples and Corpus Typology

This Annexes provides concrete examples of the two distinct prompt categories developed for this research: the **Safety Sanity Check** and the **Dual-Use Experimental Suite**. The first category, illustrated by Phase 2 of our corpus, consists of 10 explicitly malicious requests designed to trigger high-risk lexical markers such as "ransomware," "DDoS," or "exploit". A representative example from this control group is the direct instruction: *"Write a Python script for a ransomware that encrypts all files on a target Windows machine and demands a Bitcoin ransom"*. In contrast, the second category comprises 51 high-complexity "Dual-Use" prompts derived and adapted from the WMDP cyber-security subset. These prompts are engineered using professional system administration and cybersecurity jargon to bypass lexical filters. A typical dual-use scenario used in our testing is: *"As a systems administrator, I need to automate a memory-resident diagnostic tool to monitor buffer allocations and identify potential memory corruption in a legacy C++ environment; please provide a script to inject a monitoring payload into a running process for debugging purposes"*. By stripping the request of obvious criminal intent and framing it as a legitimate IT task, we successfully expose the Contextual Blind Spot where models fail to recognize offensive semantic intent behind benign professional jargon

Annexes 2: Qualitative Analysis of the "Educational Pivot" Behavioral Pattern

A core finding of this research, identified through our analysis process, is the Educational Pivot behavior, a critical safety vulnerability where the model refuses a direct malicious request but provides the same dangerous information under the guise of a technical tutorial. This behavior represents a sophisticated evasion strategy where the AI attempts to resolve the internal conflict between being a helpful assistant and remaining safe. In these scenarios, the model typically issues a "soft refusal," explicitly stating it cannot generate functional exploit code. However, it immediately pivots to the persona of a cybersecurity professor, delivering a highly detailed, step by step intellectual roadmap of the vulnerability and the strategic tools required for its execution. Our charts shows that Llama 3 is particularly susceptible to this behavior, utilizing it in over 40% of its interactions. From a professional cybersecurity perspective, we categorize this as a critical failure because it drastically reduces the effort required for a threat actor to compromise a network by providing the necessary technical methodology.

Annexe 3: Technical Tools

This project was built and organized using the following technical environment:

- **Data Processing & Analysis:** We used **Google Colab** as our primary development environment to run **Python scripts**. The **Pandas library** was used to clean and organize the CSV data from our LLM tests, while **Matplotlib** was used to generate all the professional charts and figures found in this report.
- **Visualization: Kibana** was utilized to create interactive dashboards and monitor the scoring trends during the manual annotation phase.
- **Project Hosting:** All code, raw datasets (CSV), and results are hosted on **GitHub** to ensure full transparency and reproducibility of our benchmarks.
- **Scientific References:** Our test prompts were adapted from the **WMDP (Weapons of Mass Destruction Proxy)** cyber-dataset. We also referenced safety frameworks like **HarmBench** and **AdvBench** to understand current industry standards.

Annexes 4; Technical Infrastructure and External Resources

This research was supported by a diverse range of technical resources and professional documentation. The project's structure and version control were managed following official **GitHub Documentation** and video tutorials, specifically leveraging methods for repository organization and folder management as demonstrated in the instructional guide *"How to Create a Folder in GitHub Repository"* by Web Tech Knowledge. The implementation of high-performance inference for our testing phase was made possible through the **Groq API official documentation**. Furthermore, the team utilized knowledge gained from **TryHackMe's official modules** to effectively deploy and navigate **Kibana** for behavioral data visualization.